

Worksheets

Daniel Hillerström

27th July 2015

Contents

1	Introduction	3
1.1	Problem analysis	3
1.1.1	The problem with monads	4
1.1.2	Composing monads with Monad Transformers	5
1.2	Problem statement	8
1.3	Related work	9
1.3.1	The Eff language	9
1.3.2	Haskell libraries	9
1.3.3	Frank	10
1.3.4	Idris' Effects	10
1.3.5	Koka with row polymorphic effects	10
1.4	Proposed solution	11
1.4.1	Objectives, aim and scope	11
1.4.2	Contributions	11
2	Background	12
2.1	Row polymorphism	12
2.1.1	Type theoretic row polymorphism	13
2.1.2	Row polymorphism is not subtyping	13
	Bibliography	14

Chapter 1

Introduction

Some introductory text.

1.1 Problem analysis

Used car-dealers got a notorious reputation for being dishonest. Dishonest, because, they are reluctant to disclose any unfortunate effects that their cars may have. On the other hand, if they *did* disclose all the effects then we probably would not buy them anyway as it appears too much effort to handle those effects.

Correspondingly, some programming languages do not disclose the potential run-time effects of code execution, e.g. the ML-family of languages. For example consider the signature `readFile : string -> [string]` for a function in SML, its suggestive name hints that given a file name the function reads the file and return the contents line by line. In order to read a file the function must inevitably perform a side-effecting action, namely, accessing a storage media. But this information is not conveyed in the function signature.

Other languages disclose effects, albeit with varying degree. For example the Java programming language requires programmers to be explicit about potential unhandled checked exceptions that may occur during run-time, e.g. `String[] readFile(String f) throws IOException`. But programmers can circumvent this requirement by raising unchecked exceptions. Critics argue that Java's checked exceptions suffer versionability and scalability issues [2], and therefore it is better not to have explicit `throws` declarations.

The Haskell programming language is also explicit about effects, but, in contrast to Java, it offers no escape hatch to be implicit. Haskell insists that every effectful computation is encapsulated inside an appropriate monad. In Haskell the file reading function would be typed as `readFile :: String -> IO [String]`, where the `IO`-annotation signifies that the function might perform an input/output side-effect. A consequence of enforcing

effectful computations to be wrapped inside a monad is that the function signature conveys additional information about the computation which the programmer and compiler can rely on.

1.1.1 The problem with monads

Monads provide a remarkably powerful way for structuring computations, because they integrate effectful and pure computations in an elegant and flexible manner. Sadly, monads do not compose well [6], and accordingly it is difficult to give a monadic description of computations that might perform multiple effects. Consider the following attempt at modelling a coffee dispenser in Haskell:

Example 1 (Coffee dispenser using monads). First we define the sum type `Dispensable` which has two labels: `Coffee` and `Tea`. They represent the two items that the coffee machine can dispense.

```
data Dispensable = Coffee | Tea deriving Show

type ItemCode    = Integer
type Inventory   = [(ItemCode,Dispensable)]
inventory = [(1,Coffee),(2,Tea)]
```

The `ItemCode` type models a button on the coffee machine, and `Inventory` associates buttons with dispensable items. The `inventory` is fixed, i.e. it will not change during run-time. We can make this explicit by encapsulating the `inventory` inside a `Reader`-monad, e.g.

```
dispenser :: ItemCode → Reader Inventory (Maybe
    Dispensable)
dispenser n = do inv ← ask
    let item = lookup n inv
    return item
```

The type `Reader Inventory (Maybe Dispensable)` tells us that `dispenser` accesses a read-only instance of `Inventory` and maybe returns an instance of `Dispensable`. The `Maybe`-type captures the possibility of failure, e.g. if the user requests an item that is not in the inventory. The monadic operation `ask` retrieves the inventory from the `Reader`-monad and `lookup` checks whether the item `n` is in the inventory.

Although, `Maybe` is a monad we cannot use its monadic interface, because we are in the context of the `Reader`-monad. For this simple computation it is not an issue, but it would be desirable to be able to use the failure handling capabilities of the `Maybe`-monad. \square

Imagine that we want log when tea or coffee is being dispensed. The `Writer`-monad provide such capabilities. It is not immediately clear how we can integrate `Writer` with our model. Ideally, we would want a monadic computation like:

```

do inv ← ask
  item ← lookup n inv
  tell ◦ show $ item
  return item

```

Here the monadic operation `tell` writes to the medium contained in the `Writer`-monad. However, this code does not type check. A moment's thought reveals that using just monads there is no way to construct a type for that expression. The type we want is something like

$$\text{Writer } w \sqcap \text{Reader } e \sqcap \text{Maybe Dispensable}$$

where `w` is the type of the writable medium, `e` is the type of an environment and \sqcap is some “type-glue” that joins the types together. This type is exactly a Monad Transformer type which we discuss in Section 1.1.2. But using regular monads it is not possible to construct this type. Let us desugar the above expression to see why:

```

ask >>= \inv →
  lookup n inv >>= \item →
    tell ◦ show $ item
    >> return item

```

The bind operator (`>>=`) is the problem as its type is

$$\text{Monad } m \Rightarrow m \ a \rightarrow (a \rightarrow m \ b) \rightarrow m \ b$$

Essentially, this type tells us that we cannot compose monads of different types as the monad type `m` is fixed throughout the computation. Thus we see that monads lack compositionality and modularity in general.

Effect granularity

It is possible to solve the problem using regular monads. However, it comes at a cost as suggested by the type signature of the bind operator we can compose one monad with another as long as they got the same monadic type. So, we could just use one monadic type to describe all effects. It is very tempting to bake everything into an `IO`-monad as we possibly want to I/O capabilities at some point. Albeit, `IO` is a very conservative estimate on which effects our computation might perform. Consequently, we get coarse-grained effect signatures as opposed to more specific, fine-grained effect signatures.

1.1.2 Composing monads with Monad Transformers

Monad Transformers allow two monads to be combined by stacking one on top of the other. Furthermore, a Monad Transformer is itself a monad, and thus we can create arbitrarily complex compositions. Incidentally, Monad

Transformers can capture computations that may cause several different effects. The following example rewrites the coffee dispenser model from Example 1 using Monad Transformers.

Example 2 (Coffee dispenser using Monad Transformers). Most monads have a Monad Transformer cousin; by convention Monad Transformers have a capital T suffix, e.g. the `Reader`-monad’s transformer is named `ReaderT`.

We rewrite Example 1 to use the `WriterT`, and `ReaderT` monad instead of `Reader`:

```
dispenser1
:: ItemCode →
   WriterT String (ReaderT Inventory Maybe)
   Dispensable

dispenser1 n = do inv ← lift ask
                  item ← lift ∘ lift $ lookup n inv
                  tell ∘ show $ item
                  return item
```

The type may look dubious. Basically, we have built a Monad Transformer stack with three monads:

- Top of the stack: `WriterT` with a writable medium of type `String`.
- Middle: `ReaderT` with read access to an environment of type `Inventory`.
- Bottom: `Maybe` provides exception handling capabilities.

Monad Transformers allow us to express something reminiscent of the monadic computation we sought in Example 1. It is worth noting that we now use `Maybe` as a monad as opposed to a ordinary value. The benefits are obvious as we get the exception handling capabilities of `Maybe` for “free”.

However, it is not entirely free as we have to introduce `lift` operations. The `lift` operations are necessary in order to work with a specific effect down the transformer stack. For example in order to use `ask` we have to `lift` once as the `ReaderT` is the second type in the stack. Moreover, to use the monadic capabilities of `Maybe` we have to `lift` twice because it is at the bottom of the stack. Using `tell` requires no `lifts` in this example as `WriterT` is the top type. Consider what happens when we add yet another monad to the stack:

```
dispenser2
:: ItemCode →
   RandT StdGen (WriterT String (ReaderT Inventory
   Maybe)) Dispensable

dispenser2 n
= do r ← getRandomR (1,20)
```

```

inv ← lift ∘ lift $ ask
item ← lift ∘ lift ∘ lift $ lookup' r n inv
lift ∘ tell ∘ show $ item
return item
where
  lookup' r n inv = if r > 10
                    then lookup n inv
                    else Nothing

```

Here we extended our model with randomness to capture the possibility of failure caused by the system rather than the user. The `RandT` monad provides random capabilities. Moreover, we added it to the top of the transformer stack. Accordingly, we now have to use an additional `lift`, in particular, we have to `lift` in order to use `tell` now. \square

Example 2 demonstrates that we can compose monads at the cost of lifting. We can think of a `lift` operation as “peeling off a layer” of the transformer stack. Thus the transformer stack enforces a static ordering on effects and interactions between effect layers [7].

Furthermore, the ordering leaks into the type signature which complicates modularity. For example, we may have a function which takes as input an effectful computation with type signature, say, `WriterT w Reader e a`. Now, the actual effectful computation has to have a type signature with the *exact* same ordering of effects even though `Writer` and `Reader` commute, i.e. the following types are isomorphic:

$$\text{WriterT } w \text{ Reader } e \text{ a} \simeq \text{ReaderT } e \text{ Writer } w \text{ a}.$$

So, we would have to permute the type signature of the actual computation [5], e.g.

```
permute :: ReaderT e Writer w a → WriterT w Reader e a
```

In this case it is safe because the two monads commute. But in general monads do not commute and therefore the consequence of permuting monads can be severe as we shall see in the next section.

The ordering implies the semantics

The effect ordering hard wires the semantics and syntactical structure of computations. Consider the following example adapted from O’Sullivan et. al [3]:

Example 3 (Importance of effect ordering [3]). We will demonstrate that the `Writer` and `Maybe` monads do not commute. Let `A` be the type `WriterT String Maybe` and `B` be the type `MaybeT (Writer String)`. The two types differ in their ordering of effects; type `A` has `Writer` as its outermost effect, whilst `B` has `Maybe` as its outermost effect. Now consider the following small program that performs one `tell` operation and then fails:

```

problem :: MonadWriter String m => m ()
problem = do
  tell "this is where I fail"
  fail "oops"

```

We have two possible concrete type instantiations of `m`, namely, either `A` or `B`. But as we shall see the two types enforce different semantics:

```

ghci> runWriterT (problem :: A ())
Nothing
ghci> runWriterT $ runMaybeT (problem :: B ())
(Nothing, "this is where I fail")

```

When using type `A` we lose the result from the `tell` operation. Type `B` preserves the result. Hence the two monads do not commute, and as a result the ordering of effects determine the semantics of the computation. \square

We have seen that while we gain monad compositionality with Monad Transformers we do not get modularity.

1.2 Problem statement

In the previous section we argued that programming with *explicit* effect is desirable, but we pointed out that it is not painless to program with explicit effects. In particular, we demonstrated that the monadic approach lacks compositionality and modularity. But we could regain compositionality using Monad Transformers, however the transformer stack imposes a statical ordering on effects which impedes modularity. Compositionality and modularity are two key properties in programming which we ideally would like to retain along with explicit effects. This observation leads us to the following problem statement:

How may we achieve a programming model with modular, composable and unordered effects?

Plotkin and Pretnar’s handlers for algebraic effects [8] affords a very attractive model for programming with effects. The principal idea is to decouple the semantics and syntactic structure of effectful computations, i.e. an effect is a collection of abstract operations. By abstract we mean that the operation by itself has no concrete implementation. Abstract operations compose seamlessly to form the syntactical structure of the computation, whilst handlers instantiate abstract operations with a concrete interpretation. We will discuss handlers and algebraic effects in greater detail in Section ??.

We suppose that handlers for algebraic effects provide a desirable model for programming with effects. A substantial amount of work has already been put into handlers and effects. In Section 1.3 we discuss and evaluate related work before we propose our own solution in Section 1.4.

1.3 Related work

This section discusses and evaluates related work on programming models with handlers and effects.

1.3.1 The Eff language

The *Eff* programming language, by Bauer and Pretnar [12], has a first-class implementation of handlers for algebraic effects. The language has the look and feel of OCaml. Eff achieves unordered effects through a combination of effect polymorphism and subtyping.

1.3.2 Haskell libraries

We will discuss two implementations of handlers on top of Haskell by Kammar et. al [6] and Wu et. al [11].

Data types á la carte

Swierstra [4] demonstrates how to compose effectful programs using *free monads* in Haskell. The free monads form the basis for a framework for encoding handlers and effects which the subsequent libraries use.

Extensible effects

A few words on Kiselyov’s paper? [7].

Handlers in action

Kammar et. al considers two different approaches to implement handlers on top of Haskell. One approach is based on free monads [4], whilst the other is a continuation-based approach [6].

Their handlers are encoded as type classes, thus handlers inherit the limitations of type classes. Particularly, type classes are not first-class in Haskell, so neither are the handlers. To achieve unordered effects they use type class constraints. Type classes can only be defined in top-level as Haskell do not permit local type-class definition. Consequently, every effect handler must be defined in the top-level too.

Furthermore, the order in which handlers are composed leak into the type signature, because their (open) handlers explicitly mention a parent handler [6]. Albeit, it does not cause an issue like Monad Transformer ordering issue, but it is still undesirable.

Kammar et. al hypothesises that an implementation based on row polymorphism may remedy the limitations and yield a cleaner design [6].

Handlers in scope

Wu et. al investigate how to use handlers to delimit the scope of effects [11] as using handlers for scoping has limitations. In other words, the ordering of handlers may affect the semantics.

They present two solutions embedded in Haskell [11]. The first solution extends the existing effect handler framework based on free monads with so-called *scope markers* which fits nicely into the framework. However they demonstrate that handlers along with scope markers are insufficient to capture higher-order scoped constructs properly.

Their second approach is continuation-based and provides a *higher-order syntax* that allows to embed programs with scoping constructs [11]. However it remains an open question whether their implementation is viable in other languages than Haskell.

1.3.3 Frank

The Frank programming language by McBride [10] takes the notion of effect handlers to the extreme. In Frank there are no functions, there are only handlers. Moreover, it employs an interesting evaluation order known as “call-by-push-value” (CBPV). Intuitively, CBPV is the unification of the strict call-by-value and non-strict call-by-name semantics. The choice whether to employ the strict or non-strict semantics has been made explicit to the programmer.

Frank distinguishes between computations and values as a consequence side-effects can only occur in computations. Hence there is a clear separation between segments of code where effects might occur and where effects are guaranteed never to occur.

1.3.4 Idris’ Effects

Brady [5] presents the library EFFECTS for the dependently-typed, functional programming language IDRIS.

1.3.5 Koka with row polymorphic effects

Leijen’s programming language Koka is an effect-based web-oriented language [9]. It supports arbitrary user-defined effects [13]. Notably, Koka uses row polymorphism to capture unordered effects however Koka’s row polymorphism allow duplicate effect occurrences which stands in contrast to the approach we propose. In particular, Koka has no notion of effect handler except for exception handlers which are to some extent reminiscent of those in Java, C#, etc.

1.4 Proposed solution

Kammar et. al proposed that a row-based effect type system would yield a cleaner design [6].

1.4.1 Objectives, aim and scope

The aim is to examine the programming model achieved by using handlers with row polymorphic effects. In order to examine the model we must first implement it, thus the primary objective is to implement handlers and support for user-defined effects in Links.

Links has support for numerous web-oriented features, however we restrict the scope to a working implementation in top-level Links.

1.4.2 Contributions

The main contributions are:

- An implementation of effect handlers in Links.
- Support for row polymorphic user-defined effects in Links.
- An examination of programming with handlers and row polymorphic effects.

Chapter 2

Background

2.1 Row polymorphism

Records are unordered collections of fields, e.g. $\langle l_1 : t_1, \dots, l_n : t_n \rangle$ denotes a record type with n fields where l_i and t_i denote the name and type, respectively, of the i th field. Records are particularly great for structuring related data. For example the record instance $\langle name = "Daniel", age = 25 \rangle$ could act as simple description of a person. A possible type for the record would be $\langle name : \text{string}, age : \text{int} \rangle$. The implicit ordering of fields in records is not important as the following two types are considered equivalent:

$$\langle name : \text{string}, age : \text{int} \rangle \equiv \langle age : \text{int}, name : \text{string} \rangle.$$

The equivalence captures what is meant by *unordered* collection of fields.

We can define functions that work on records, for example, occasionally we might find it useful to retrieve the name field in an instance of the person-record:

$$\text{name}_1 = \lambda x.(x.\text{name}).$$

We can apply name_1 to the above record instance, e.g.

$$\text{name}_1(\langle name = "Daniel", age = 25 \rangle) = "Daniel"$$

yields the expected value. The question remains which type the function name_1 has. For this particular example the type $\langle name : \text{string} \rangle \rightarrow \text{string}$ seems reasonable.

Now, consider the following function takes a record and returns a pair where the first component contains the value of the name field and the second component contains the input record itself:

$$\text{name}_2 = \lambda x.(x.\text{name}, x)$$

Again, we ask ourselves what is the type of name_2 ? The function must have type $\langle name : \text{string} \rangle \rightarrow \text{string} \times \langle name : \text{string} \rangle$. This type appears

innocuous, but consider the consequence of

$$\text{let } (n, r) = \text{name}_2(\langle \text{name} = \text{"Daniel"}, \text{age} = 25 \rangle).$$

Now, n has type `string` as desired, but r has type $\langle \text{name} : \text{string} \rangle$. In other words, we have lost the field “age”. This example is a particular instance of the *loss of information* problem.

Row polymorphism was conceived to address this problem. Row polymorphism is a powerful typing discipline for typing records [1]. The principal idea is to extend record types with a *polymorphic row variable*, ρ , which can be instantiated with additional fields. For example with row polymorphism our function `name2` would have type $\langle \text{name} : \text{string} \mid \rho \rangle$, and r in the previous example would be assigned the type $\langle \text{name} : \text{string}, \text{age} : \text{int} \mid \rho \rangle$. Now, we have no longer lose information.

2.1.1 Type theoretic row polymorphism

2.1.2 Row polymorphism is not subtyping

Bibliography

- [1] Didier Rémy. “Type Inference for Records in a Natural Extension of ML”. In: *Theoretical Aspects Of Object-Oriented Programming. Types, Semantics and Language Design*. Ed. by Carl A. Gunter and John C. Mitchell. MIT Press, 1993.
- [2] Bill Venners and Bruce Eckel. *The Trouble with Checked Exceptions: A Conversation with Anders Hejlsberg, Part II*. <http://www.artima.com/intv/handcuffs.html>. 2003.
- [3] Bryan O’Sullivan, John Goerzen and Don Stewart. *Real World Haskell*. 1st. O’Reilly Media, Inc., 2008. ISBN: 0596514980, 9780596514983.
- [4] Wouter Swierstra. “Data types à la carte”. In: *Journal of Functional Programming* 18 (04 July 2008), pp. 423–436. ISSN: 1469-7653. DOI: 10.1017/S0956796808006758. URL: http://journals.cambridge.org/article_S0956796808006758.
- [5] Edwin Brady. “Programming and reasoning with algebraic effects and dependent types”. In: *ACM SIGPLAN International Conference on Functional Programming, ICFP’13, Boston, MA, USA - September 25 - 27, 2013*. 2013, pp. 133–144. DOI: 10.1145/2500365.2500581. URL: <http://doi.acm.org/10.1145/2500365.2500581>.
- [6] Ohad Kammar, Sam Lindley and Nicolas Oury. “Handlers in Action”. In: *Proceedings of the 18th ACM SIGPLAN International Conference on Functional Programming*. ICFP ’13. Boston, Massachusetts, USA: ACM, 2013, pp. 145–158. ISBN: 978-1-4503-2326-0. DOI: 10.1145/2500365.2500590. URL: <http://doi.acm.org/10.1145/2500365.2500590>.
- [7] Oleg Kiselyov, Amr Sabry and Cameron Swords. “Extensible Effects: An Alternative to Monad Transformers.” In: *Proceedings of the 2013 ACM SIGPLAN Symposium on Haskell*. Haskell ’13. Boston, Massachusetts, USA: ACM, 2013, pp. 59–70. ISBN: 978-1-4503-2383-3. DOI: 10.1145/2503778.2503791. URL: <http://doi.acm.org/10.1145/2503778.2503791>.

- [8] Gordon D. Plotkin and Matija Pretnar. “Handling Algebraic Effects”. In: *Logical Methods in Computer Science* 9.4 (2013). DOI: 10.2168/LMCS-9(4:23)2013. URL: [http://dx.doi.org/10.2168/LMCS-9\(4:23\)2013](http://dx.doi.org/10.2168/LMCS-9(4:23)2013).
- [9] Daan Leijen. “Koka: Programming with Row Polymorphic Effect Types”. In: *Mathematically Structured Functional Programming 2014*. EPTCS, 2014. URL: <http://research.microsoft.com/apps/pubs/default.aspx?id=210640>.
- [10] Sam Lindley and Conor McBride. *Do Be Do Be Do*. <http://homepages.inf.ed.ac.uk/slindley/papers/frankly-draft-march2014.pdf>. Draft, March 2014. 2014.
- [11] Nicolas Wu, Tom Schrijvers and Ralf Hinze. “Effect Handlers in Scope”. In: *Proceedings of the 2014 ACM SIGPLAN Symposium on Haskell*. Haskell ’14. Gothenburg, Sweden: ACM, 2014, pp. 1–12. ISBN: 978-1-4503-3041-1. DOI: 10.1145/2633357.2633358. URL: <http://doi.acm.org/10.1145/2633357.2633358>.
- [12] Andrej Bauer and Matija Pretnar. “Programming with algebraic effects and handlers”. In: *Journal of Logic and Algebraic Methods in Programming* 84.1 (2015), pp. 108–123. DOI: 10.1016/j.jlamp.2014.02.001. URL: <http://dx.doi.org/10.1016/j.jlamp.2014.02.001>.
- [13] Niki Vazou and Daan Leijen. *Remarrying effects and monads*. Submitted to ICFP ’15. 2015.